

You can't spell "devshell"

without
"hell"

Today's lesson

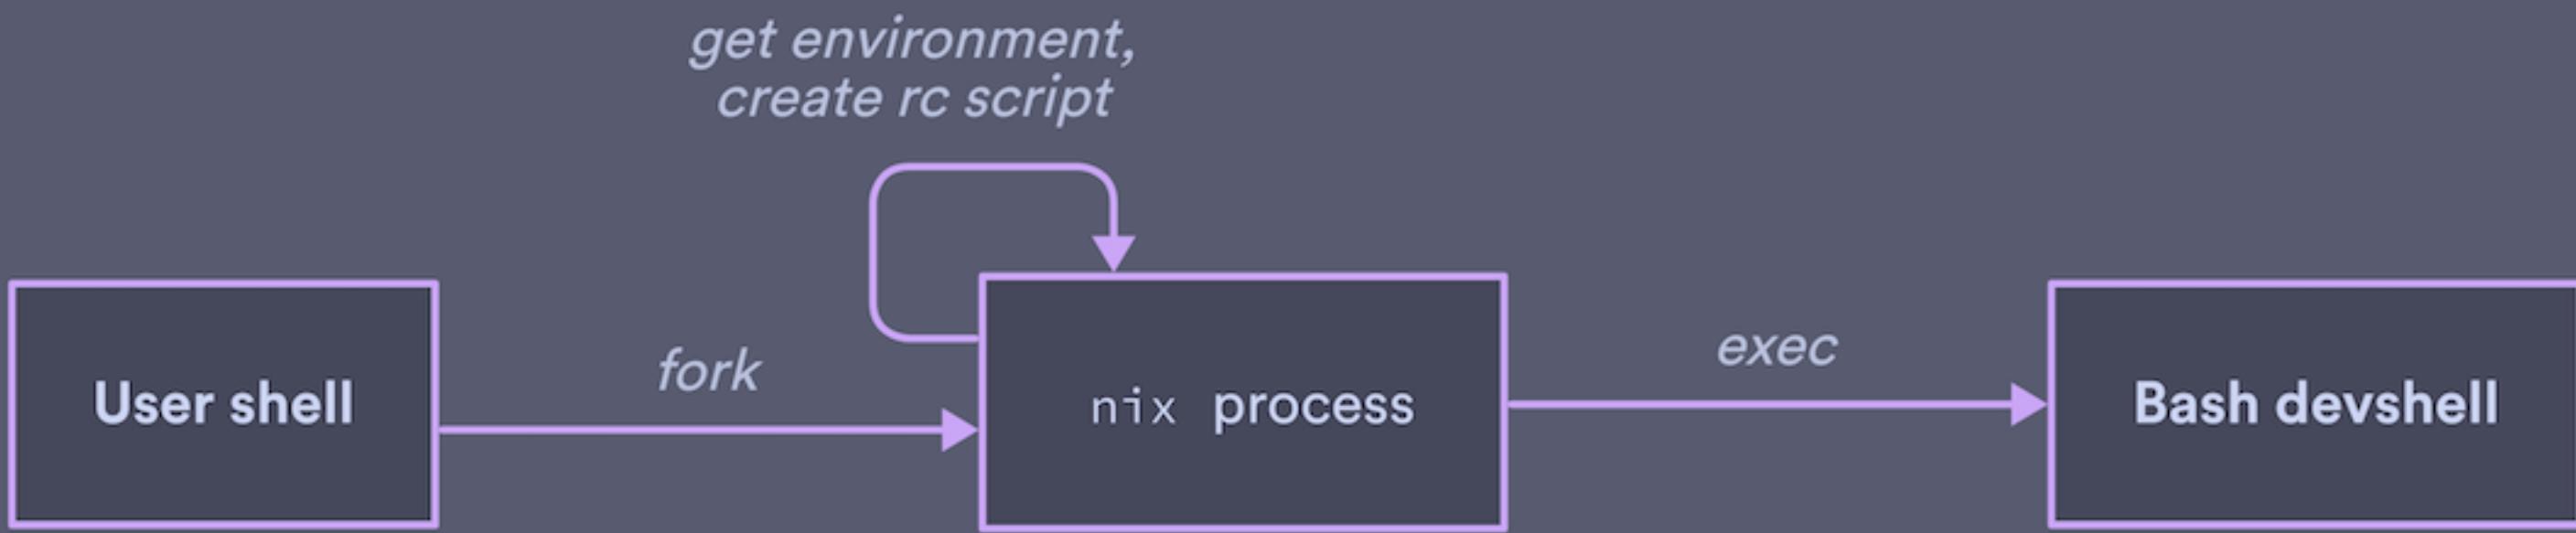
- How does nix develop work?
- What if we want more?
- flox activate
- Why fish is objectively the best shell

The environment

- Filesystem + environment = observable universe
- Control them both = control what's observable
- Strategy
 - Deterministically set the environment
 - Manipulate filesystem via symlinks

nix
develop

nix develop flow



If you want to see for yourself

In develop.cc

- CmdDevelop::run
- getBuildEnvironment
- getDerivationEnvironment
- struct BuildEnvironment
- makeRcScript
- getShellOutputpath

Trivia!

- `nix-shell` and `nix develop`
 - Designed for debugging builds
 - Both recreate the `build environment` of a derivation
 - Not the same as a general purpose developer environment

Try running `buildPhase` inside `nix develop`

Trivia!

```
$ cd flox
$ nix develop
$ buildPhase
no Makefile or custom buildPhase, doing nothing
!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
$cargoBuildLog is either undefined or does not point to a valid file location!
By default the installFromCargoBuildLogHook will expect that cargo's output
was captured and can be used to determine which binaries should be installed
(instead of just guessing based on what is present in cargo's target directory)
```

If you are defining your own custom build step, you have two options:

1. Set `doNotPostBuildInstallCargoBinaries = true;` and ensure the installation steps are handled as appropriate.
2. ensure that cargo's build log is captured in a file and point \$cargoBuildLog at it

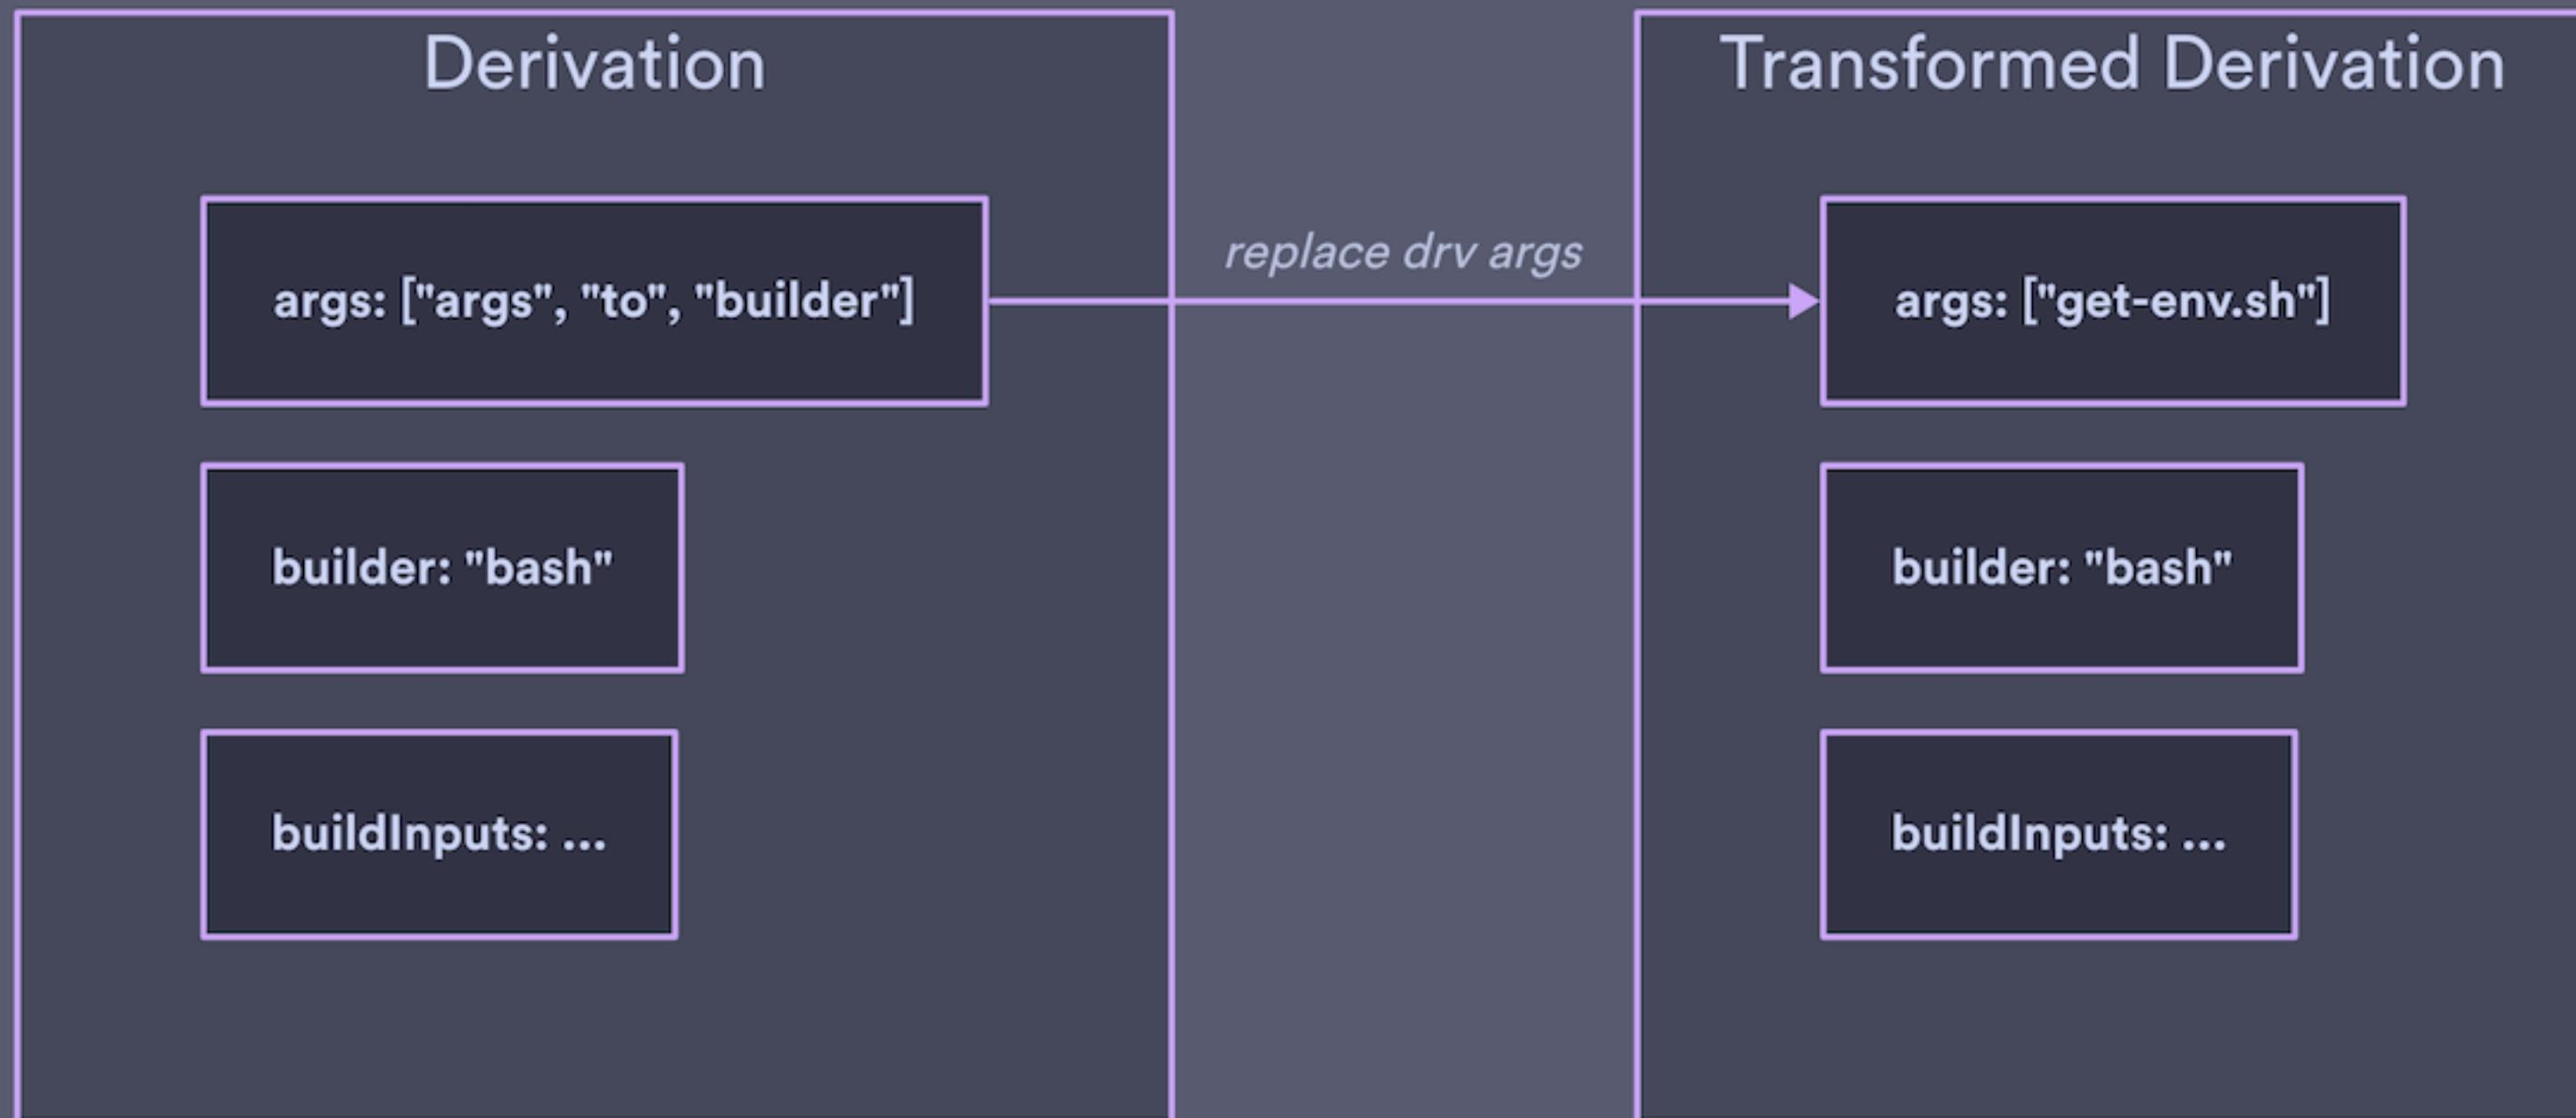
At a minimum, the latter option can be achieved with a build phase that runs:

```
cargoBuildLog=$(mktemp cargoBuildLogXXXX.json)
cargo build --release --message-format json-render-diagnostics >"$cargoBuildLog"
!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
```

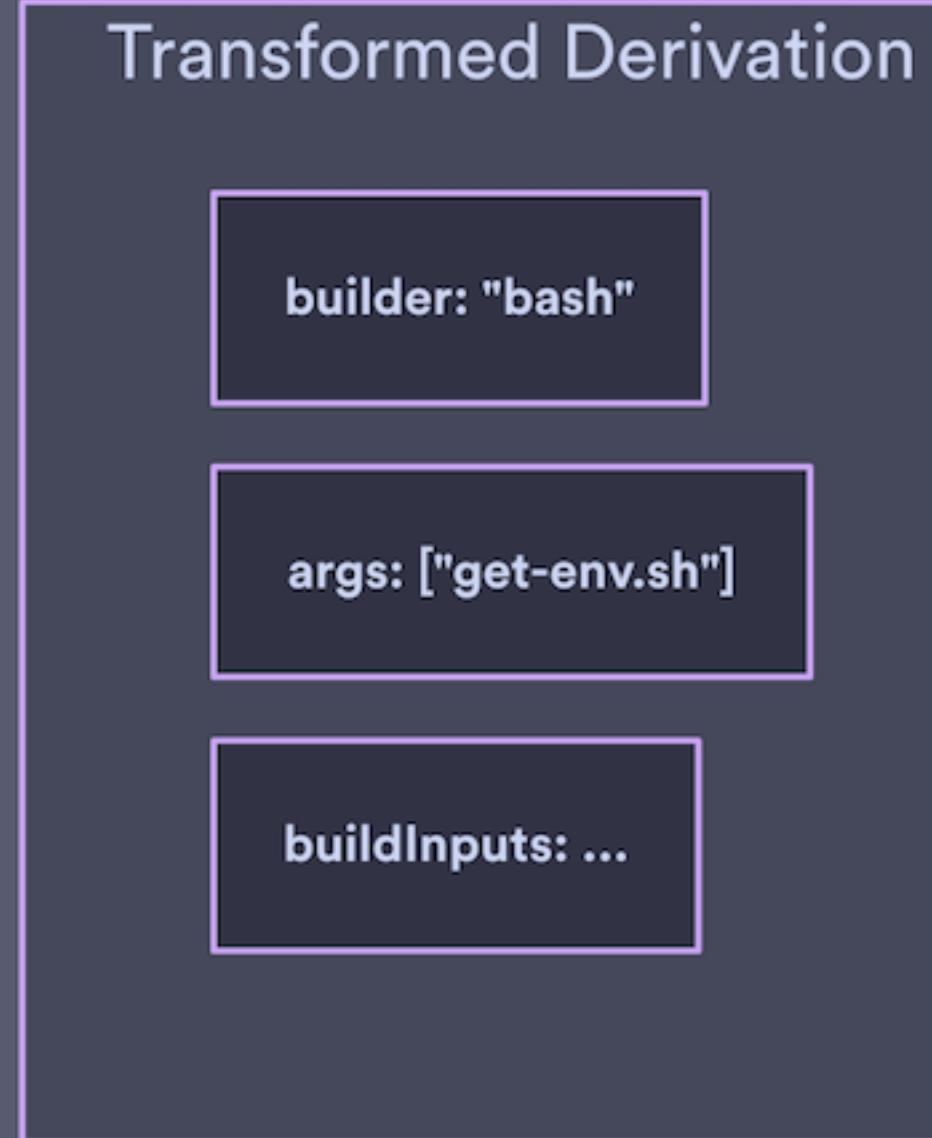
First steps

- Run `nix develop .#foo`
- Find the derivation for `foo`

Transform Derivation



Get Environment JSON



get-env.sh

```
# ...

# This populates PATH, etc from the derivation inputs
if [[ -n $stdenv ]]; then
    source $stdenv/setup
fi

# ...

__vars=$(declare -p)
__functions=$(declare -F)

__dumpEnv() {
    # Construct JSON object *line by line*
    # from __vars and __functions
}

# ...
```

Getting the function definitions

```
if ! [[ $__line =~ ^declare\ -f\ (.*) ]]; then continue; fi
__fun_name="${BASH_REMATCH[1]}"
__fun_body=$(type $__fun_name)
if [[ $__fun_body =~ \{(.*)\} ]]; then
    if [[ -z $__first ]]; then printf ',\n'; else __first=; fi
    __fun_body=${BASH_REMATCH[1]}
    printf "      "
    __escapeString "$__fun_name"
    printf ':'
    __escapeString "$__fun_body"
else
    printf "Cannot parse definition of function '%s'.\n" "$__fun_name" >&2
    return 1
fi
```



Zach Mitchell - Flox

Getting the variable definitions

```
if ! [[ $__line =~ ^declare\ (-[^ ])\ ([^=]* ) ]]; then continue; fi
local type="${BASH_REMATCH[1]}"
local __var_name="${BASH_REMATCH[2]}"

if [[ $__var_name =~ ^BASH_ || \
      $__var_name =~ ^COMP_ || \
      # <snip>
    ]]; then continue; fi

if [[ -z $__first ]]; then printf ',\n'; else __first=; fi

printf "    "
__escapeString "$__var_name"
printf ': {' 

# FIXME: handle -i, -r, -n.
if [[ $type == -x ]]; then
    printf '"type": "exported", "value": '
    __escapeString "${!__var_name}"
elif [[ $type == -- ]]; then
    printf '"type": "var", "value": '
    __escapeString "${!__var_name}"
elif [[ $type == -a ]]; then
    printf '"type": "array", "value": ['
    local __first2=1
    # <snip>
    printf '['
elif [[ $type == -A ]]; then
    printf '"type": "associative", "value": {\n'
    local __first2=1
    declare -n __var_name2="$__var_name"
    # <snip>
    printf '\n    }'
else
    printf '"type": "unknown"'
fi

printf "}"
```

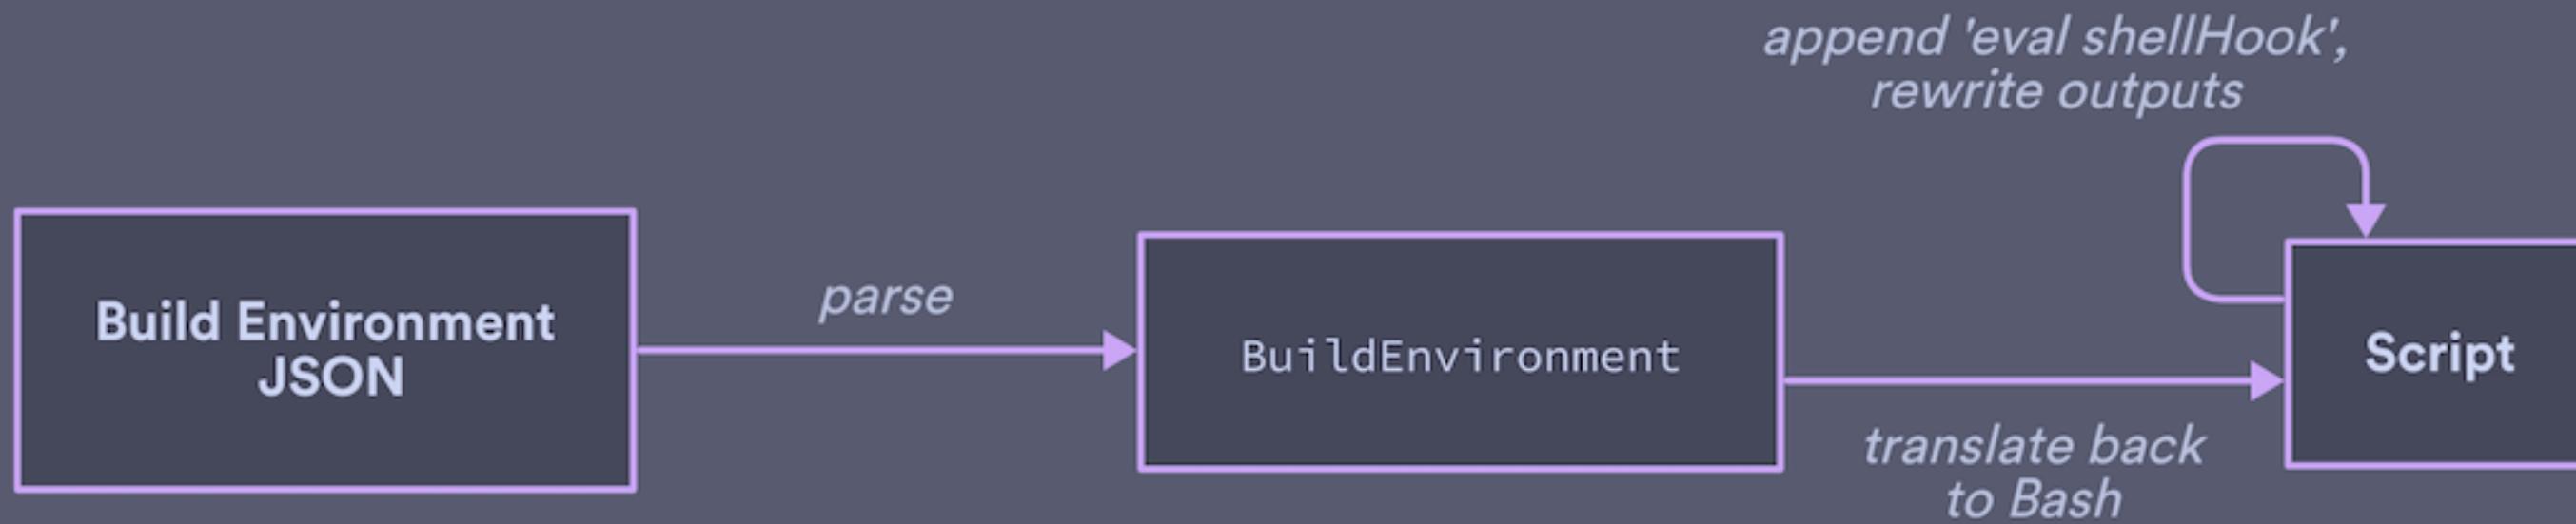


Viewing the JSON

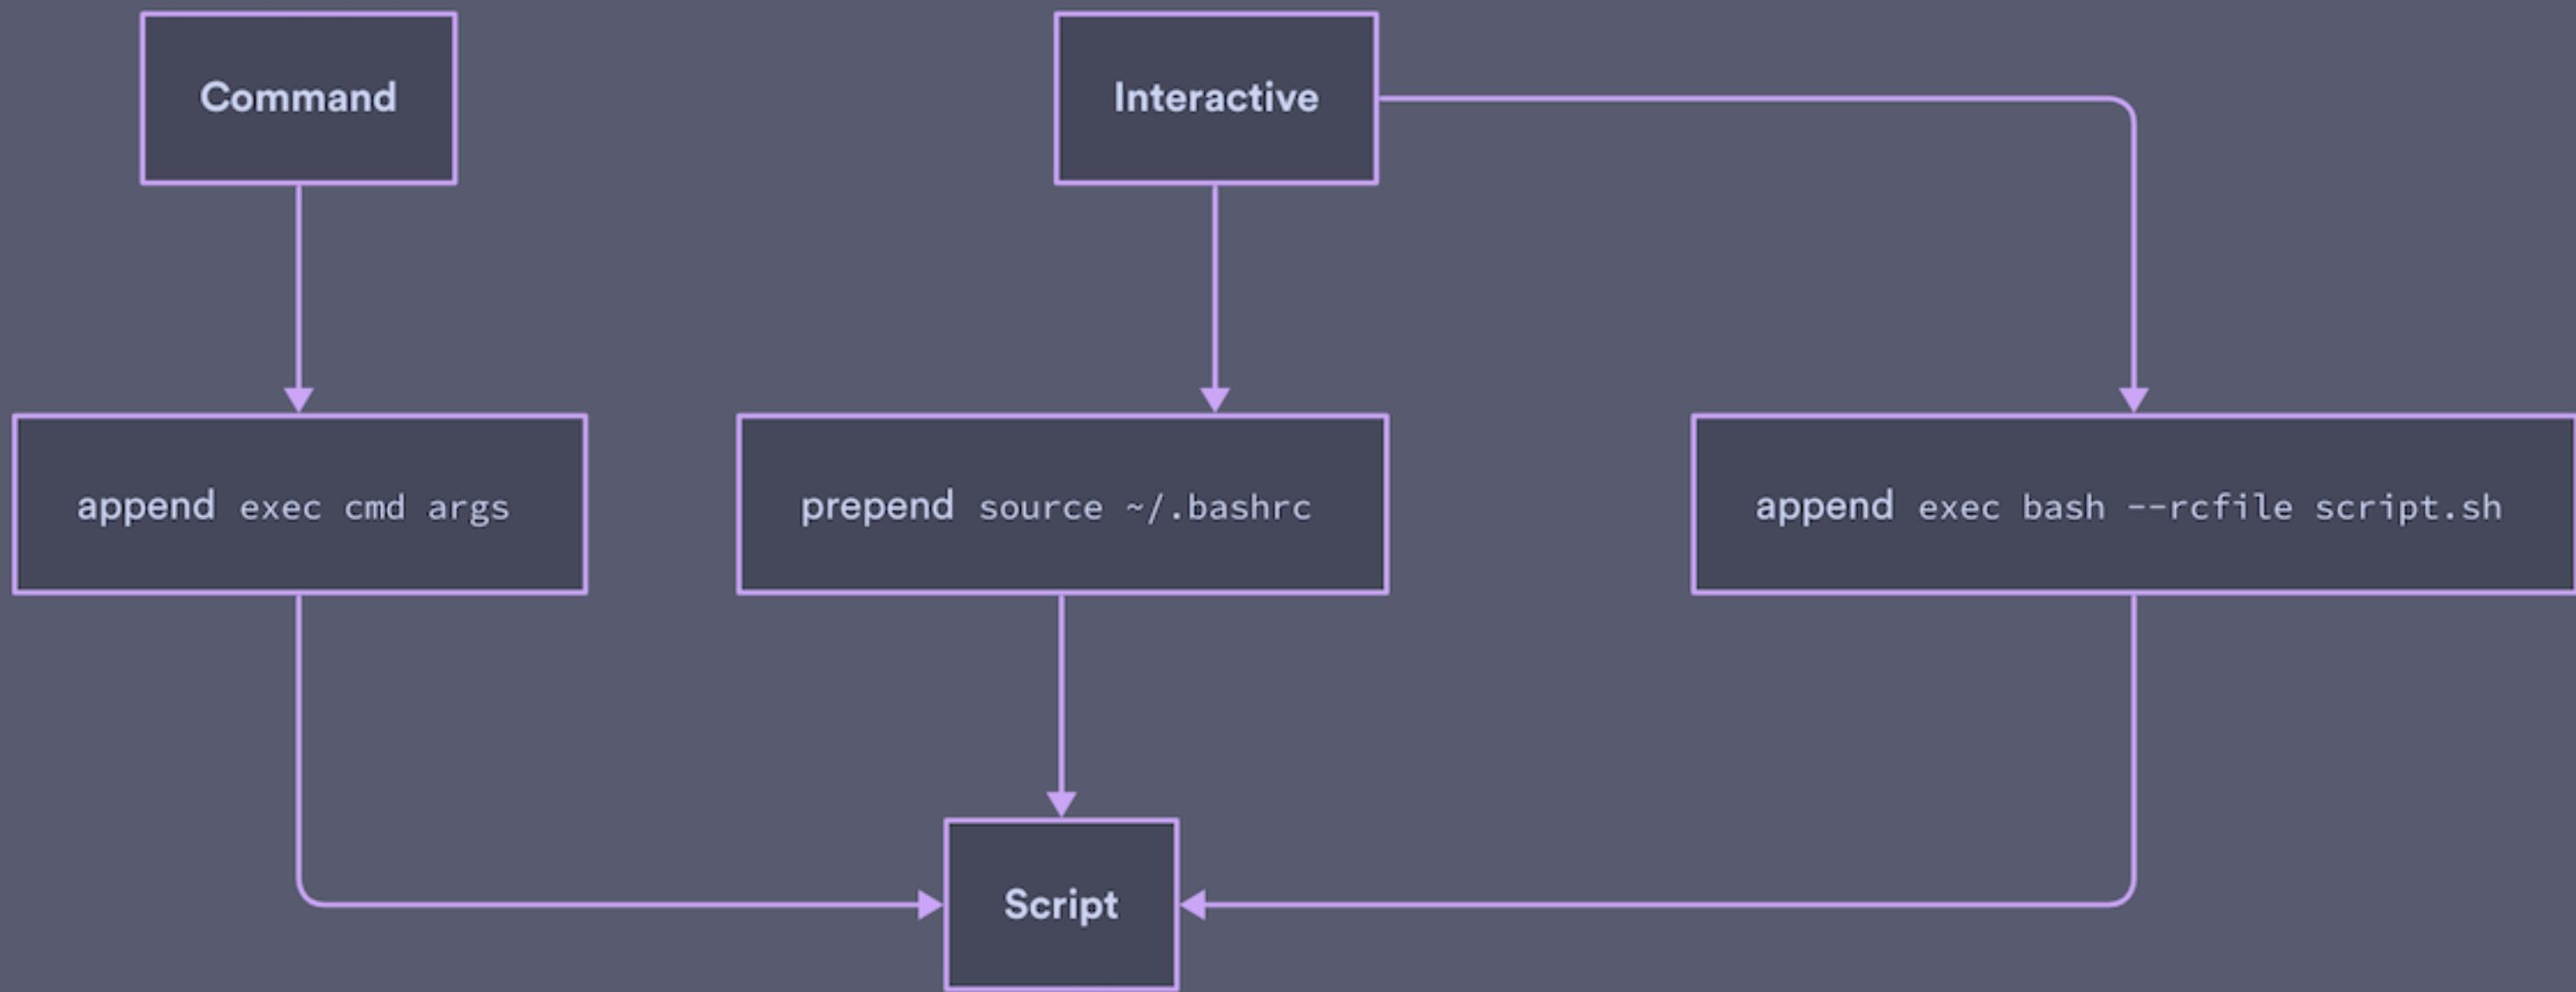
```
$ nix print-dev-env --json
```

```
{
  "bashFunctions": {
    "myFunc": "...body..."
  },
  "variables": {
    "CC": {
      "type": "exported",
      "value": "/nix/store/..."
    },
  }
}
```

Make RC Script



nix develop mode



"Just exec your shell"

- `nix develop` only supports Bash
- You want to use `<shell>`
- Advised to `nix develop -c exec <shell>`
- WRONG
 - Shell's RC files will run **after** Nix setup
 - All kinds of shenanigans can happen

Modify, don't overwrite some variables

- PATH, XDG_DATA_DIRS
- Only including things from the devshell would cause major breakage
- Prepend to these instead of overwriting

Temp directories

- Create a temporary directory `nix-shell.XXXXXXX`
- Point all temp-like directories at it
 - `TMP`, `TEMP`, `TMPDIR`, `TEMPDIR`

```
$ echo $TMP  
/tmp/nix-shell.NUgD1Q
```

Rewrite outputs

- Normally \$out points at the destination in the store
- Rewrite occurrences of that store path to ./outputs/<output>
- Prevents you from writing garbage to the store during development

```
$ echo $out  
/Users/zmitchell/src/flox/main/outputs/out
```

Garbage collection

- Set `NIX_GCROOT` to script path in `nix` process
- Nix garbage collector uses `/proc` to scan for GC roots in running processes
 - See `LocalStore::findRuntimeRoots`

Write out the shell script

```
$ nix print-dev-env  
unset shellHook  
PATH=${PATH:-}  
nix_saved_PATH="$PATH"  
...  
...
```

Write out the shell script

```
$ nix build --print-out-paths --no-link \
.#devShells.aarch64-darwin.default
/nix/store/zkh8q1m320573982wvqq68wgcaf50af6-flox-dev

$ cat /nix/store/zkh8q1m320573982wvqq68wgcaf50af6-flox-dev
unset shellHook
PATH=${PATH:-}
nix_saved_PATH="$PATH"
...
```

**flox
activate**

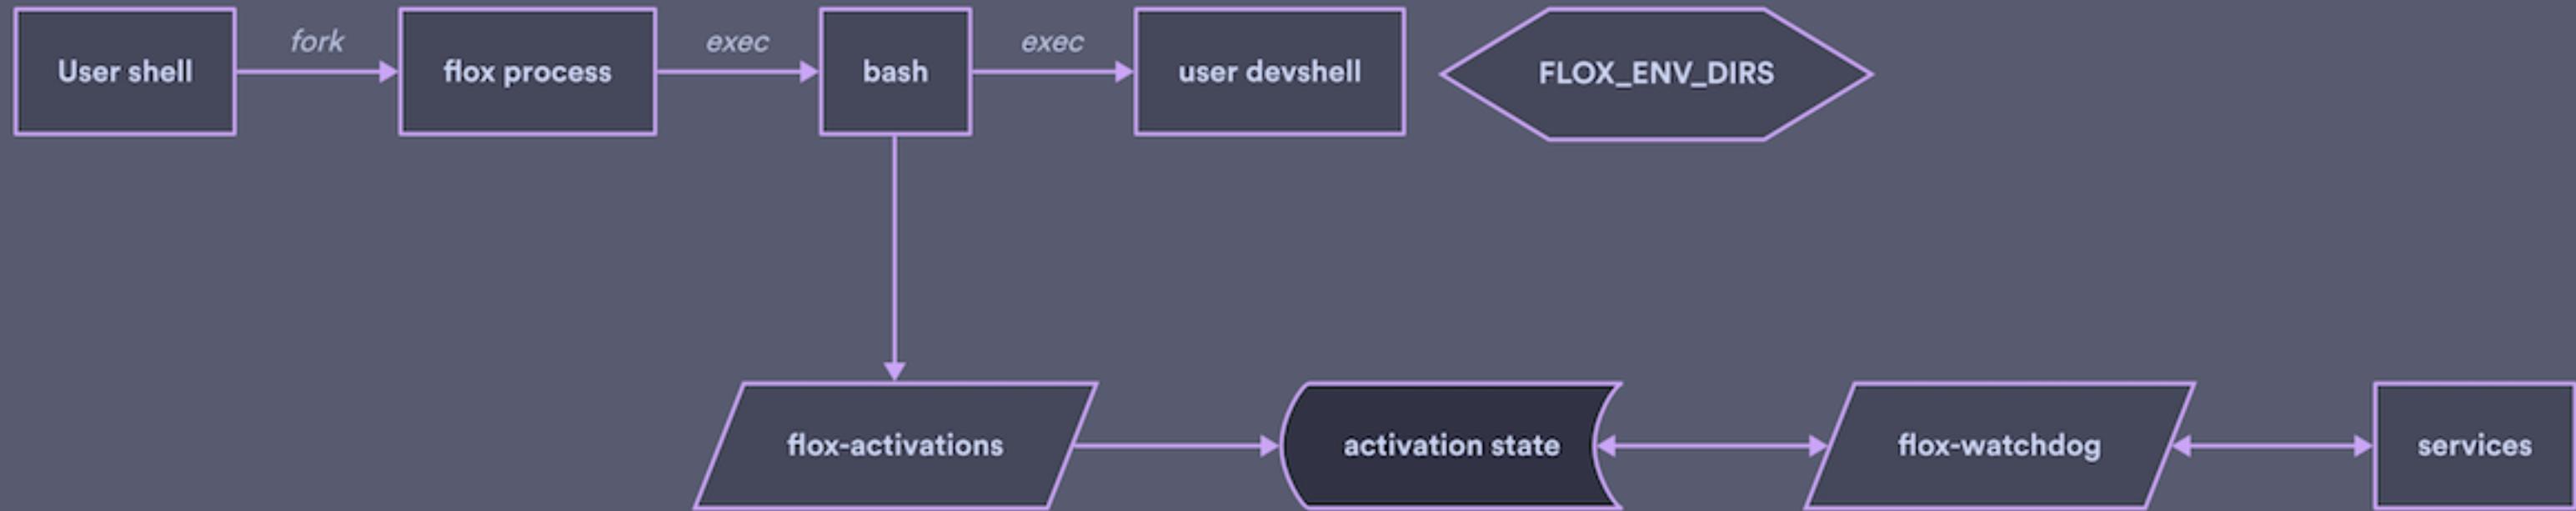
Different goals

- Same reproducibility as nix develop
- 4 different shells
 - Bash, Zsh, Fish, and tcsh
- 2 different closures
 - "develop" and "runtime" modes
- "Immediate" activation

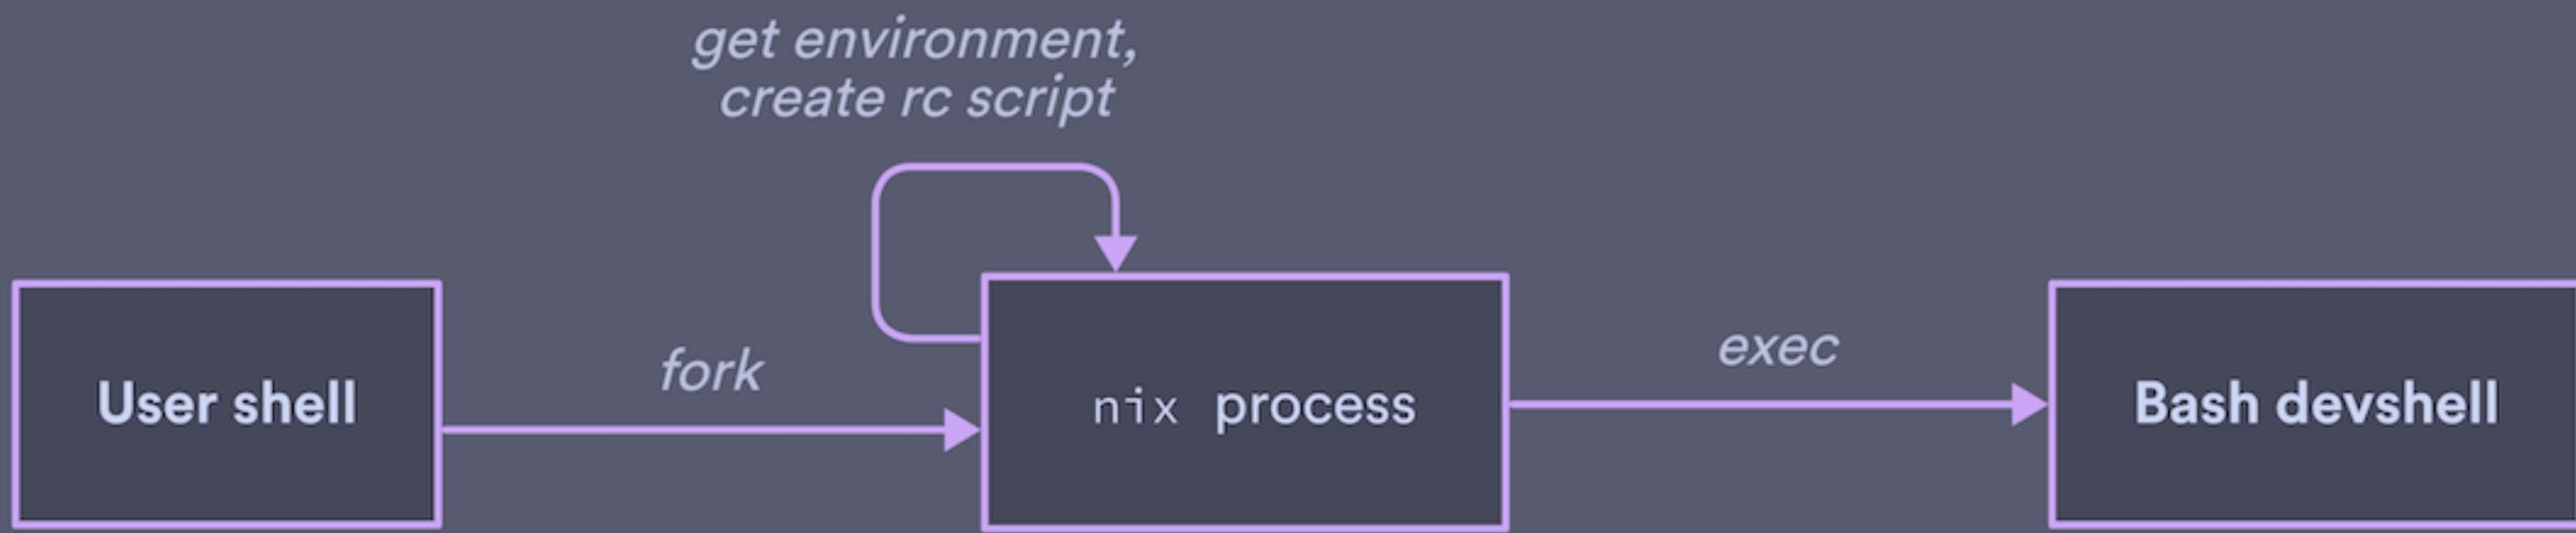
Additional requirements

- Layered activations
- Share resources between activations
- Services w/ automatic shutdown on exit
- Tidy PATH

flox activate flow



nix develop flow



Completely different architecture

- Everything needed to run the shell is already built
- Eagerly rebuilt when `flox` commands modify the environment
- Rebuilt on the fly when a `flox` command detects manual edits.

Built ahead of time

```
> tre -a .flox -l 2
.flox
├── env.json
├── .gitignore
└── env
    ├── manifest.toml # ← config file
    └── manifest.lock # ← lockfile
└── run
    ├── x86_64-linux.rfcp.run # ← symlink to "runtime" closure
    └── x86_64-linux.rfcp.dev # ← symlink to "develop" closure
└── log
└── cache
└── .gitattributes
```

Recursively merged FHS

```
› tre -a .flox/run/x86_64-linux.rfcp.dev -l 1
.flox/run/x86_64-linux.rfcp.dev
├── requisites.txt
├── etc
├── bin          # ← merged '/bin' from entire closure
├── activate     # ← 'activate' script
├── activate.d   # ← activation helpers
├── lib
├── share
├── nixpkgs-url
└── include
└── manifest.lock
```

"Immediate" activation

```
› time flox activate -- /run/current-system/sw/bin/true
```

```
-----  
Executed in 56.05 millis fish external  
usr time 31.85 millis 211.00 micros 31.64 millis  
sys time 26.39 millis 141.00 micros 26.25 millis
```

```
› time fish -c /run/current-system/sw/bin/true
```

```
-----  
Executed in 15.99 millis fish external  
usr time 7.93 millis 292.00 micros 7.64 millis  
sys time 8.01 millis 142.00 micros 7.87 millis
```

"Immediate" activation

- No evaluation on the way to activation
- Rate limiting step is often your own dotfiles
- "Quick enough" for the moment, but there's low hanging fruit

I wrote a profiler: proctrace

- Uses bpftrace to track forks/execs and timing
- Generates Gantt chart syntax for Mermaid.js
- Broken/wrong in various ways at the moment, don't use it
- <https://github.com/zmitchell/proctrace>
- proctrace.xyz

Process Trace



Completely different architecture

- Two different shell hooks
 - `hook.on-activate`: shell-agnostic setup
 - `profile.*`: sourced by user shell
- Start vs. attach
 - Start does full activation (first time)
 - Attach replays aliases + env (subsequent activations)

manifest.toml

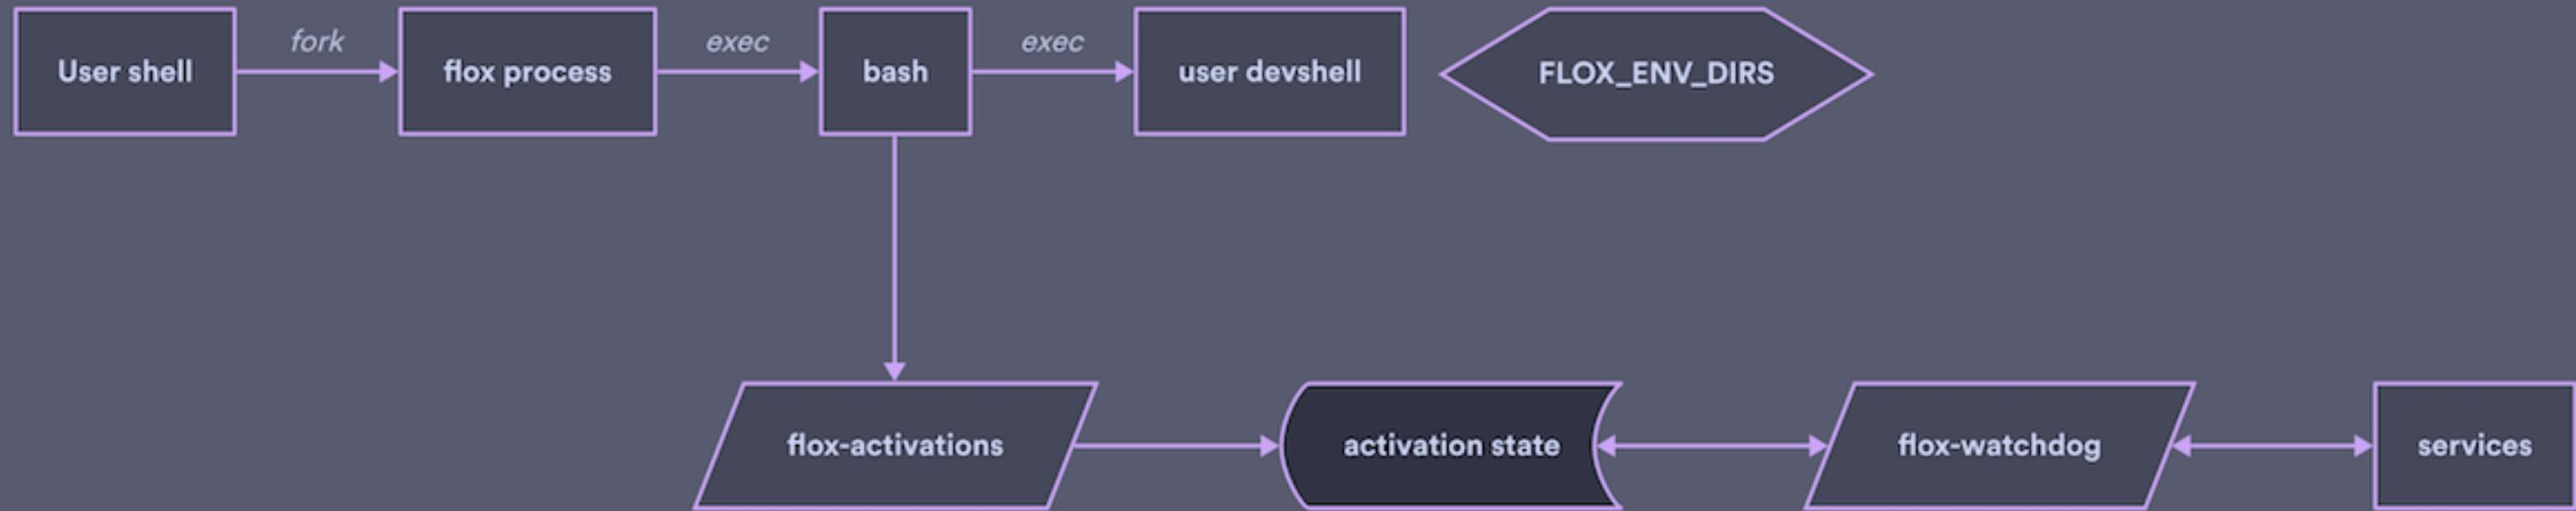
```
version = 1

[install]
ripgrep.pkg-path = "ripgrep"

[hook]
on-activate = '''
    echo "hello from hook.on-activate"
'''

[profile]
fish = '''
    echo "hello from profile.fish"
    alias myalias "my alias cmds"
'''
```

flox activate flow



Completely different architecture

- Need to keep track of instances of activated envs
 - **flox-watchdog** → background process, one per env
 - activation state → used for attach
 - One set of services shared between activations of an env
- Always possible that we're about to layer activations
 - **\$FLOX_ENV** → current activation
 - **\$FLOX_ENV_DIRS** → all activations in this shell

Generic setup

- activate script is written in Bash
- One consistent language for all of the shell-agnostic parts
- Defer shell-specific parts to the very end

First step: set FLOX_ENV_DIRS, PATH, MANPATH

- All activate code is aware that multiple activations may already be stacked
- ex.) PATH needs to contain:
 - bin/sbin directories from all active environments.
 - In most-recently-activated order
 - But also what was already in PATH

Big, difficult problem

- Implement consistent behavior
- N shell dialects with different:
 - Quoting rules
 - Array support
 - Rules for source and eval
 - etc



The Abyss

The Abyss

- 📄 Shell idiosyncrasies
- 📄 Python
- 📄 That command, it does not do what you think it does

The Abyss

- Bash and Zsh support hieroglyphics
 - "\${@@Q}"
 - "\${foo##*/}"
- Woe unto you if you forget to quote something

The Abyss

- tcsh has no equivalent of "\${foo:-}"
 - (default to empty string if unset or null)

The Abyss

- Don't tell tcsh to eval a string containing #
 - Often won't treat it as a comment character
 - '#: Command not found'
 - Other times comments out the rest of the input
 - \n and ; don't help

Trust me, just don't

The Abyss

- tcsh sometimes just...won't let you quote something

```
eval ``cmd "arg"``  
#          ^^^^^^ straight to jail
```

- Use `arg:q` to quote `arg`

The Abyss

- jk that doesn't work in backticks
- Use :Q

The Abyss

- `jk` that only works in "new" tcsh versions



Zach Mitchell - Flox

The Abyss

- tcsh doesn't have an IFS variable
 - Good luck reading PATH into an array

The Abyss

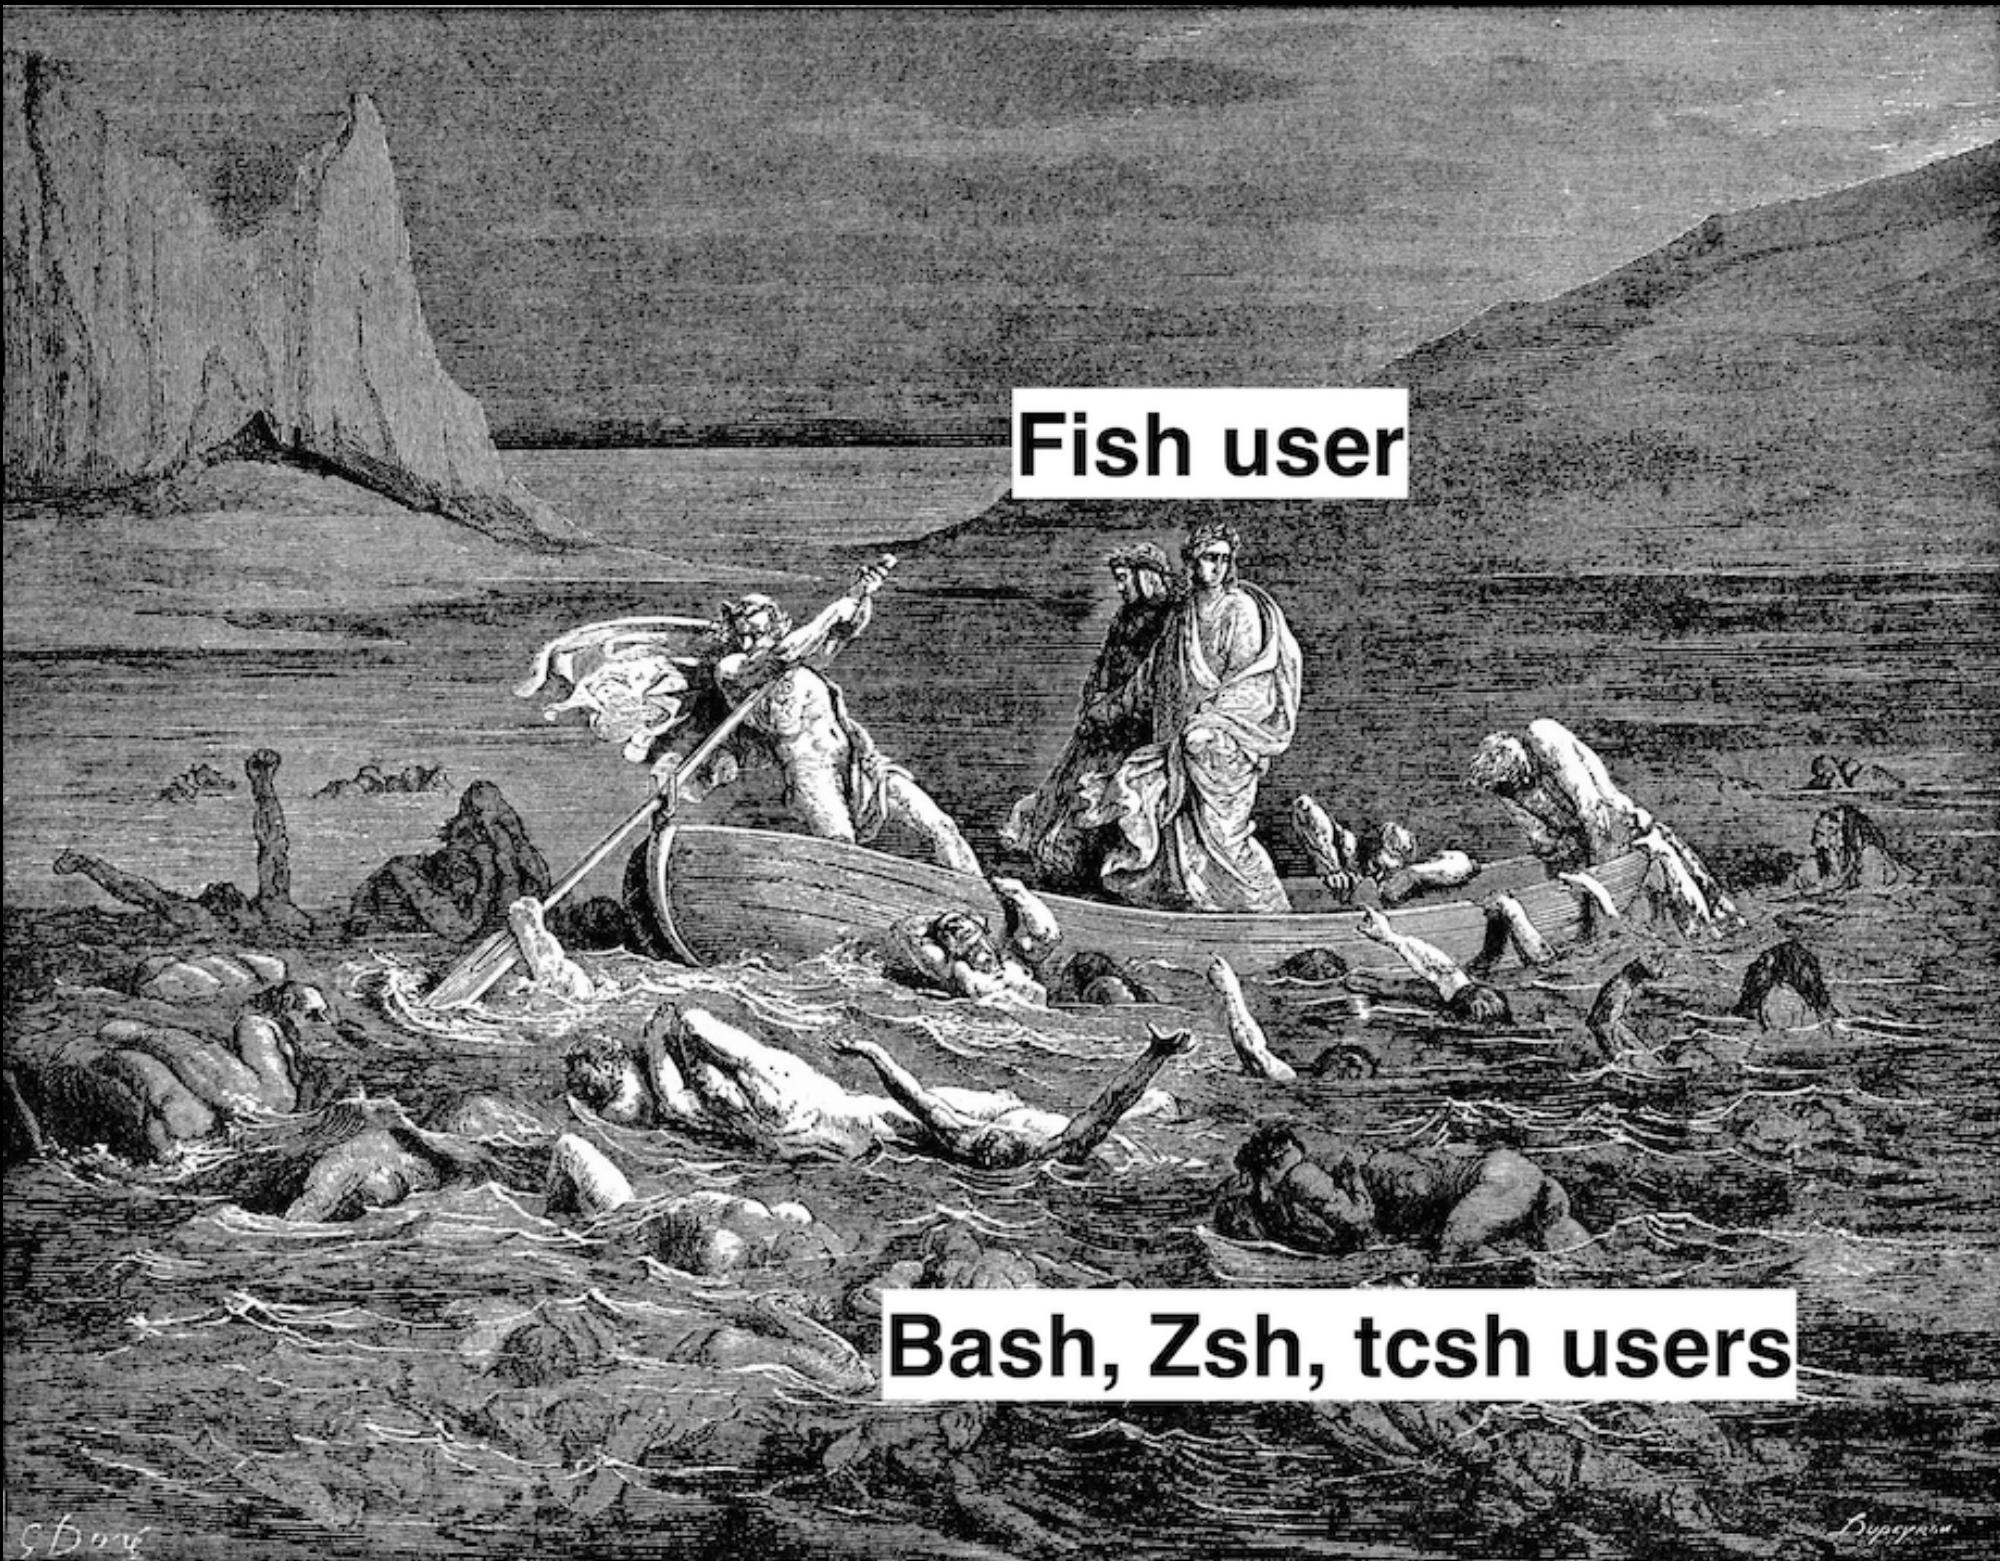
- Zsh has both `$fpath` and `$FPATH`
 - `$fpath` is an array
 - `$FPATH` is like `$PATH`
 - Set `$fpath`, all good
 - Set `$FPATH`, subshells won't load default paths

The Abyss

- Everything is a list/array in Fish
- As a result, these are printed differently
 - \$PATH
 - "\$PATH"
- fish_prompt is a function
 - Produces the format string when executed

The Abyss

- Startup files...
 - Zsh has 8
 - Bash has ~6
 - Fish has 2
 - I didn't count for tcsh

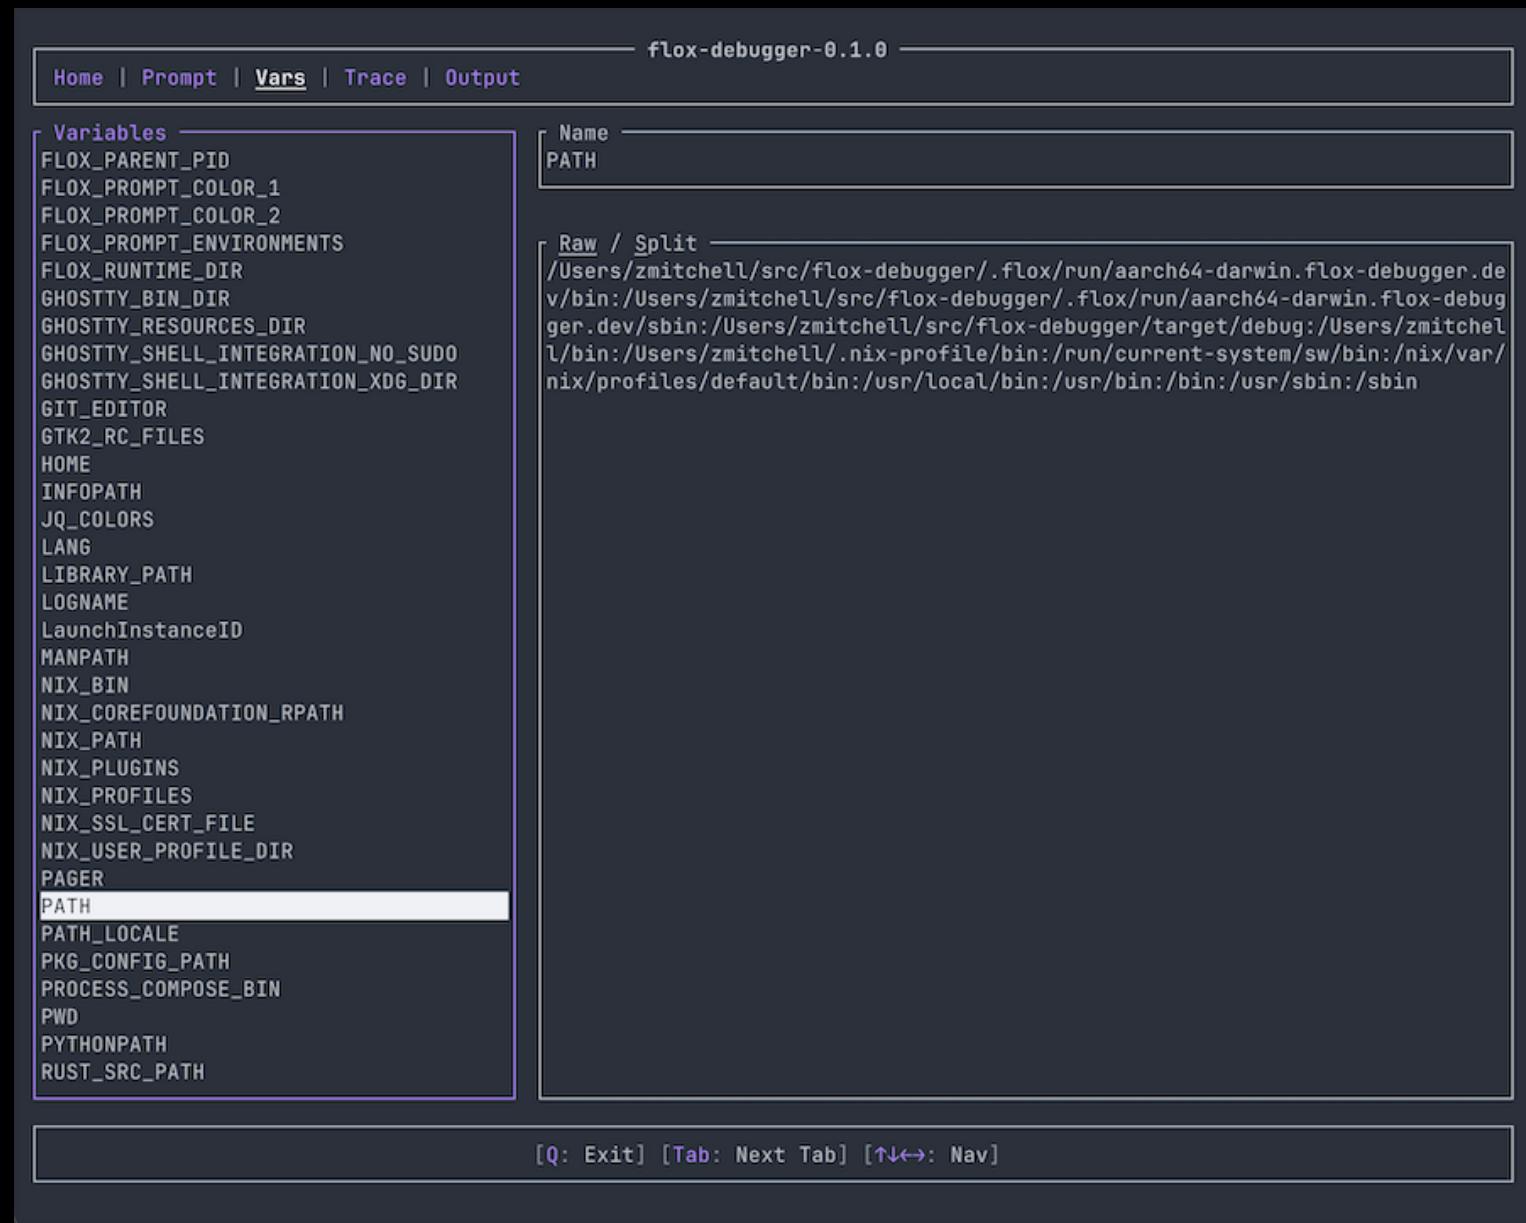




I wrote a tool for that too



I wrote a tool for that too



I wrote a tool for that too

The screenshot shows the interface of the `flox-debugger-0.1.0` tool. At the top, there is a navigation bar with links: Home, Prompt, Vars (which is underlined, indicating it's the active tab), Trace, and Output.

The main area is divided into two sections: "Variables" on the left and "Raw / Split" on the right.

Variables:

- FLOX_PARENT_PID
- FLOX_PROMPT_COLOR_1
- FLOX_PROMPT_COLOR_2
- FLOX_PROMPT_ENVIRONMENTS
- FLOX_RUNTIME_DIR
- GHOSTTY_BIN_DIR
- GHOSTTY_RESOURCES_DIR
- GHOSTTY_SHELL_INTEGRATION_NO_SUDO
- GHOSTTY_SHELL_INTEGRATION_XDG_DIR
- GIT_EDITOR
- GTK2_RC_FILES
- HOME
- INFOPATH
- JQ_COLORS
- LANG
- LIBRARY_PATH
- LOGNAME
- LaunchInstanceID
- MANPATH
- NIX_BIN
- NIX_COREFOUNDATION_RPATH
- NIX_PATH
- NIX_PLUGINS
- NIX_PROFILES
- NIX_SSL_CERT_FILE
- NIX_USER_PROFILE_DIR
- PAGER
- PATH** (highlighted in white)
- PATH_LOCALE
- PKG_CONFIG_PATH
- PROCESS_COMPOSE_BIN
- PWD
- PYTHONPATH
- RUST_SRC_PATH

Raw / Split:

Name: PATH

Raw / Split:

```
/Users/zmitchell/src/flox-debugger/.flox/run/aarch64-darwin.flox-debugger.de  
/Users/zmitchell/src/flox-debugger/.flox/run/aarch64-darwin.flox-debugger.de  
/Users/zmitchell/src/flox-debugger/target/debug  
/Users/zmitchell/bin  
/Users/zmitchell/.nix-profile/bin  
/run/current-system/sw/bin  
/nix/var/nix/profiles/default/bin  
/usr/local/bin  
/usr/bin  
/bin  
/usr/sbin  
/sbin
```

Selected:

```
/Users/zmitchell/src/flox-debugger/.flox/run/aarch64-darwin.flox-debugger.de  
v/bin
```

At the bottom, there is a footer with keyboard shortcuts: [Q: Exit] [Tab: Next Tab] [$\uparrow\downarrow\leftarrow\rightarrow$: Nav].

I wrote a tool for that too

The screenshot shows the flox-debugger-0.1.0 application window. At the top, there's a navigation bar with links: Home, Prompt, Vars, Trace (which is underlined to indicate it's active), and Output. Below the navigation bar, a message says "Current tracepoint: otherfunc3".

On the left, a "Call Stack" panel lists five frames:

- Frame #0
- Frame #1 (highlighted)
- Frame #2
- Frame #3
- Frame #4

To the right of the call stack is a "Call Site Info" panel showing the following details:

- File: /Users/zmitchell/src/flox-debugger/run.sh
- Line: 22
- Function: otherfunc3

Below these panels is a large "Call Site" panel containing the source code for the current function:

```
otherfunc() {
    fdb_traceotherfunc
    functioecho "running otherfunc"
    otherfunc2 output
}      put=""

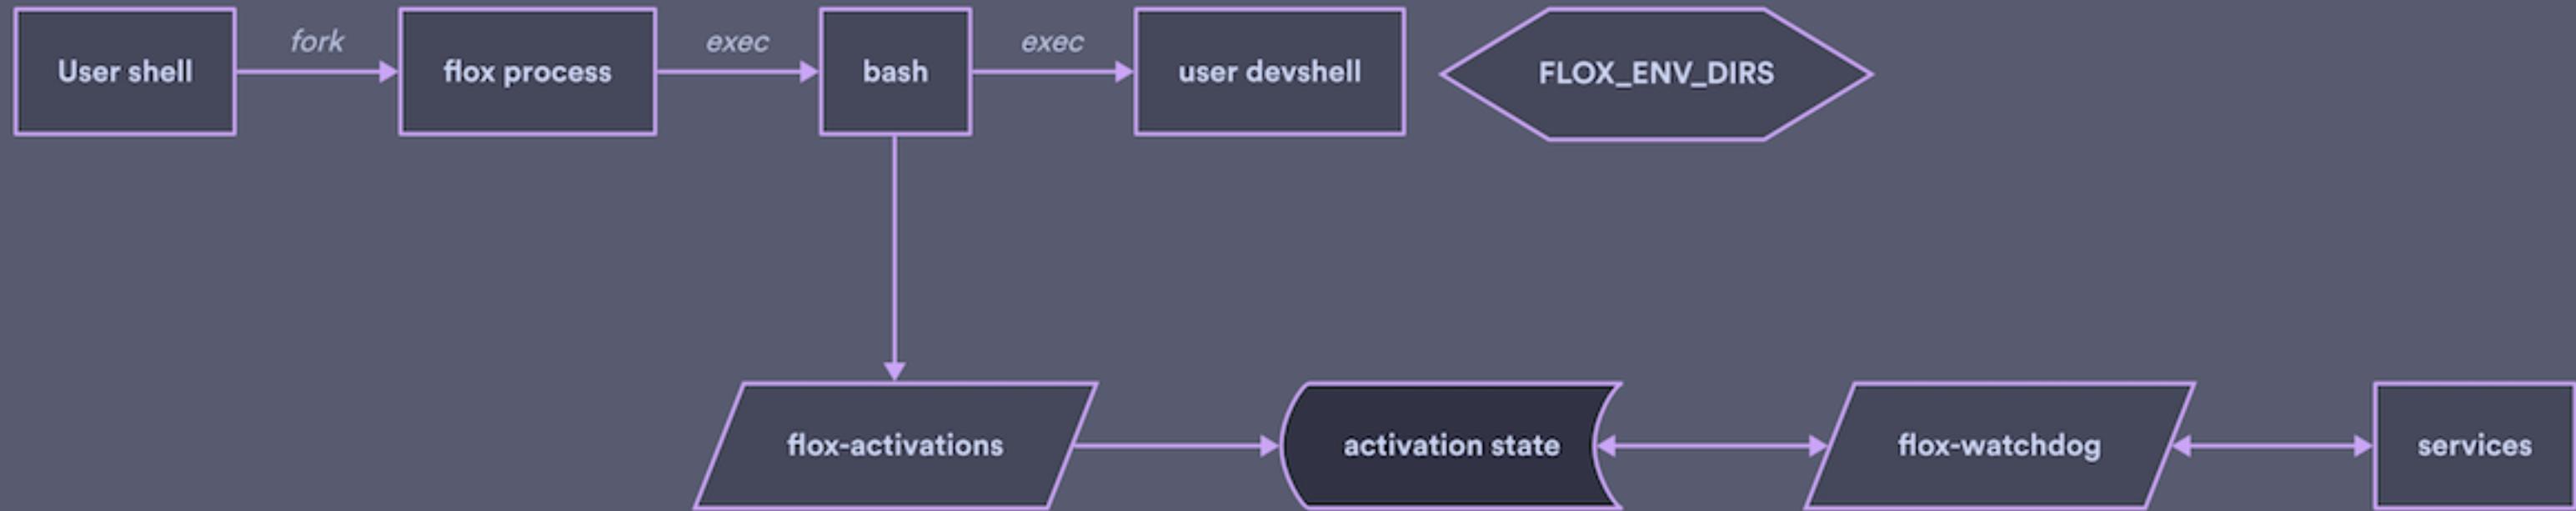
otherfulc2() { frames
    fdb_tracepoint otherfunc2E[@]}
    echo "running otherfunc2"
    otherfunc3      rame=1
}

otherfunc3() {
    fdb_tracepoint otherfunc3
    echo "running otherfunc3"
}          f a

fdb_tracepoint start
ech "Starting"local
myfunction           _SOURCE[$i]}"
fdb_tracepoint end
```

At the bottom of the window, there's a footer bar with keyboard shortcuts: [Q: Exit] [Tab: Next Tab] [: Nav]

flox activate flow



Big, difficult problem

- Implement consistent behavior
- N shell dialects with different:
 - Quoting semantics
 - Array support
 - Rules for source and eval
 - etc

Tractable problem

- Implement business logic in few languages
 - Bash for most things
 - Rust where multiple dialects is painful
- Rust
 - Mostly reduces to emitting 'export foo=bar'

ex.) Setting PATH and MANPATH

```
source <("$_flox_activations" set-env-dirs --shell bash ...)  
source <("$_flox_activations" fix-paths --shell bash ...)
```

flox-activations acts like a pure function, which makes testing and debugging easy

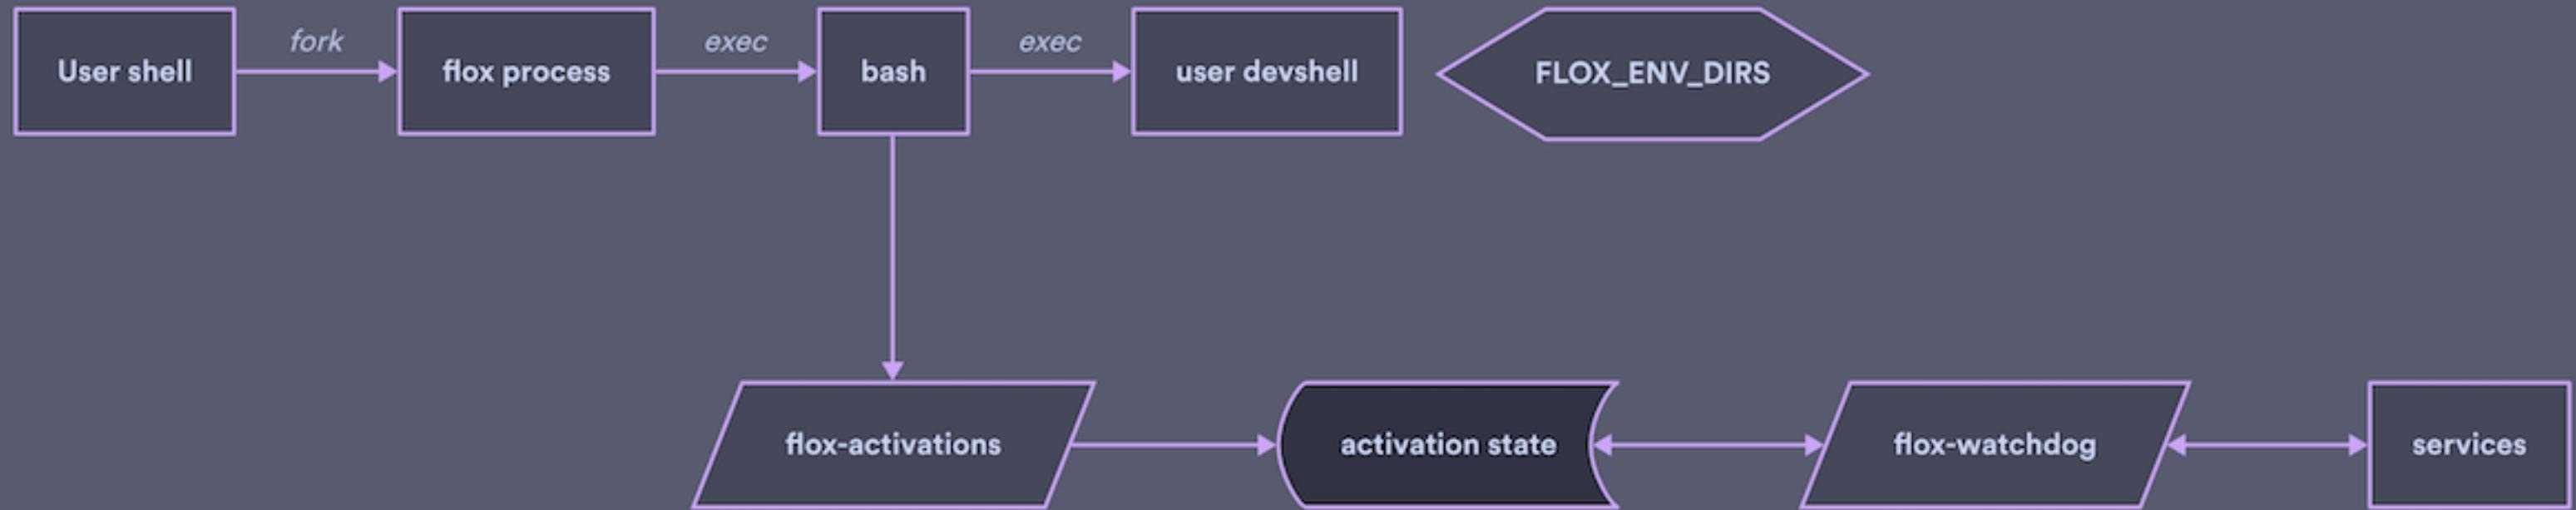
ex.) Debugging

```
› flox-activations fix-paths --shell bash \  
--env-dirs "my_env" \  
--path "foo_bin:bar_bin" \  
--manpath ""
```

Emits

```
export PATH="my_env/bin:my_env/sbin:foo_bin:bar_bin";  
export MANPATH="my_env/share/man:";
```

flox activate flow



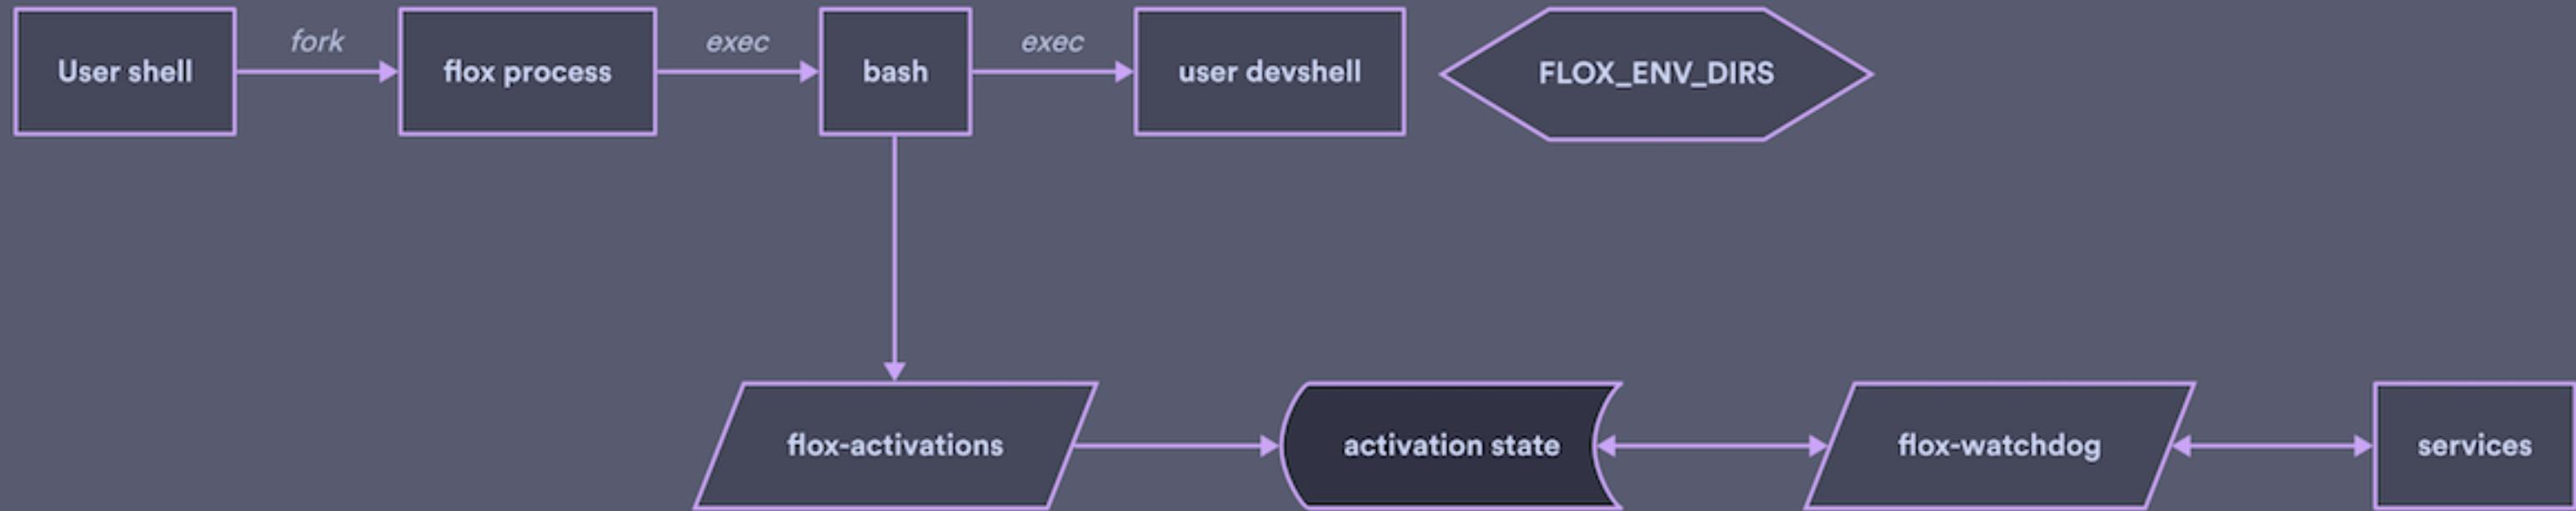
Start or attach

- `hook.on-activate` may do something expensive
- Take snapshots of the shell environment during activation
- Create diffs of those snapshots
 - What was added/deleted
- Subsequent activations apply those diffs
- `profile.*` run for every activation to set your aliases

Start services if necessary

- Done idempotently
- flox-watchdog
 - Listen for termination of last activation
 - Clean up services
 - Magic! (at the expense of gray hair)

flox activate flow



Listening for terminations

- `kill -0`
 - Doesn't differentiate between running and zombies
 - Polling is a race condition
- `tail -f /dev/null --pid <pid>`
 - Often just uses `kill -0` internally
- Read `/proc`
 - Race condition if polled



Zach Mitchell - Flox

Solution

- Platform specific APIs
- `prctl(2)` on Linux
 - `PR_SET_CHILD_SUBREAPER`
 - Configure a signal to be delivered when child exits
- `kqueue` on macOS
 - Poll a file descriptor for configurable events

Prepare RC file

- Flox **must** be the last thing that runs
- Prepare RC file that:
 - Sources user RC file
 - Runs Flox stuff (like `profile.*` scripts)
- **Very** annoying for Zsh
 - Is it interactive/login?

Links

- Flox
 - <https://flox.dev>
 - <https://github.com/flox/flox>
- Me
 - <https://github.com/zmitchell>
 - <https://tinkering.xyz>
 - @z-mitchell.bsky.social